

Event Sourcing



Dino Esposito

AUTHOR

@despos

www.software2cents.wordpress.com

In the real world you observe events.
In software, you tend to write models.



Key Points

Events

Events as **Data
Source**

Build **Projections**
of Event Data

CQRS



Models to persist



Events to log



Deep impact on system architecture

Events are also for the common application.



You have a CRUD system.



Need to know the state of the system at a given time

You may or may not be able to extract just that information with canonical and common techniques.



Newton's tree, Botanic Gardens, Cambridge


It's not that you don't
need events.

**You just don't need
events yet.**

Hello!

I'm an invoice.
Unfortunately, I
do include sales
taxes.

VAT %



X	1-1-14
Y	1-7-14
?	???

Event Sourcing

It's about ensuring that all changes made to the application state during the entire lifetime of the application are stored as a sequence of events.



This is not how the vast majority of applications work today.

Most applications work by storing the current state and using stored states to process business transactions.



Structural representation



List of ordered goods

Payment information

Shipping information



Event representation

-  Add item #1
-  Add item #2
-  Add payment info
-  Update item #2
-  Remove item #1
-  Add shipping info



Key Facts of Event Sourcing

An event is something that has happened in the past

Events are expression of the ubiquitous language

Events are not imperative and are named using past tense verbs

Have a persistent store for events

Append-only, no delete

Replay the (related) events to get to the last known state of an entity

- Replay from the beginning or a known point (**snapshot**)

An Event Is Something That Happened in the Past

- **Once stored, events are immutable**
 - Can be duplicated and replicated (for scalability reasons)
- **Any behavior associated with the event has been performed**
 - Replaying the event doesn't require to repeat the behavior
- **You don't miss a thing**
 - Track everything that happened at the time it happened
 - Regardless of the effects it produced

Data saved at a lower abstraction level

CQRS is not for just a few apps either.

CQRS and Events

CQRS is the dawn of a new software design experience. **Events** are what you find when the dawn has actually turned into a brand new day.



**Not really
revolutionary**

Events are a revolutionary new
approach.

**Not really
new**

Events are not for just a few apps either.







Power to Events



- 1. Store current state**
2. Use events to log relevant facts

- 1. Store events**
2. Build relevant snapshots of facts

Event stream

-  Q=1, Product=123
Tennis racket
-  Q=1, Product=456
Ball pack
-  CC 123456789
-  Q=2, Product=456
-  Remove Product=123
-  1 Unknown Place



Ship to:

1 Unknown Place

2	Ball pack ID=123	5.50 EUR
---	---------------------	-------------

11
EUR




Credit card:

123456789

Is it important to track **what** was added and then removed?

Is it important to track **when** an item was removed from cart?



Foundation of **Persistence**

CREATE

UPDATE

DELETE

QUERIES

CREATE Operations

Order request



Append new record

Information

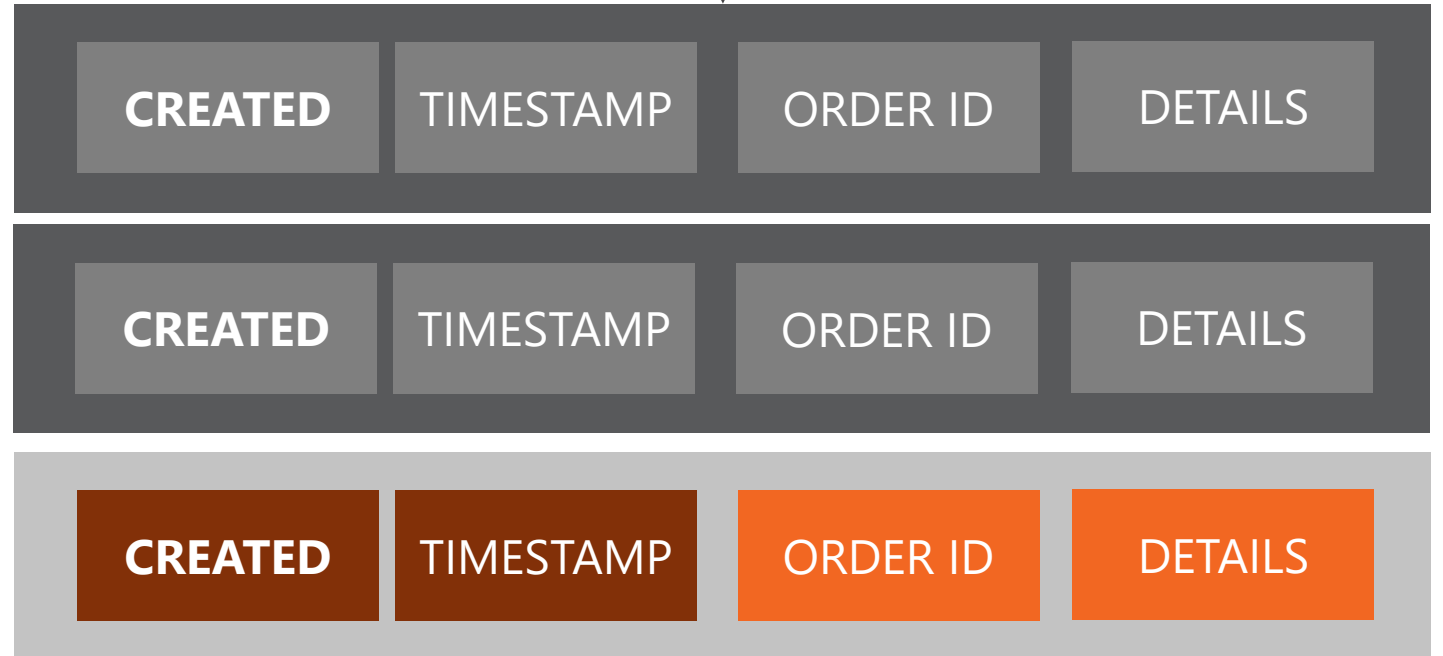
Global event ID

Code for the operation

Timestamp

Entity identifier

Full details



Transparent storage: relational, NoSQL, graph

UPDATE Operations

UPDATE request



Append new record with **DELTA**

Information

Global event ID

Code for the operation

Timestamp

Entity identifier

Changes applied

CREATED	TIMESTAMP	888998	DETAILS
CREATED	TIMESTAMP	123456	DETAILS
UPDATED	TIMESTAMP	123456	DELTA

Transparent storage: relational, NoSQL, graph

In some cases, you might want to consider storing **the full state of the entity** along with the specific event information

Event record	Current state of the entity
--------------	-----------------------------

DELETE Operations

DELETE request



Append new record to mark **logical deletion**

Information

Global event ID

Code for the operation

Timestamp

Entity identifier

Optionally reasons of deletion

CREATED	TIMESTAMP	888998	DETAILS
UPDATED	TIMESTAMP	123456	DELTA
DELETED	TIMESTAMP	123456	REASON

Transparent storage: relational, NoSQL, graph

Physical record deletion of
events in case of UNDO
functionality?

DO NOT DELETE in the
middle of the stream


Queries?

QUERY Operations

Grab all events by entity ID

64	CREATED	TIMESTAMP	123456	DETAILS
92	UPDATE D	TIMESTAMP	123456	DELTA

Replay through all events



REPLAY of Events

Get the full or partial stream of events

RavenDB query example

```
public IEnumerable<GenericEventWrapper> GetEventStream(String id)
{
    return DocumentSession
        .Query<GenericEventWrapper>()
        .Where(t => t.AggregateId == id)
        .OrderBy(t => t.Timestamp)
        .ToList();
}
```

**or an analogous
relational query**

REPLAY of Events

Generic event wrapper class

```
public class GenericEventWrapper
{
    public string EventId { get; set; }
    public string EventOperationCode { get; set; }
    public DateTime Timestamp { get; set; }
    public string AggregateId { get; set; }
    public DomainEvent Data { get; set; }
}
```

REPLAY of Events

Rebuilding the state

```
public static Aggregate PlayEvents(String id, IEnumerable<DomainEvent> events)
{
    var aggregate = new Aggregate(id);
    foreach (var e in events)
    {
        if (e is AggregateCreatedEvent)
            aggregate.Create(e.Data);

        if (e is AggregateUpdatedvent)
            aggregate.Update(e.Data);

        if (e is AggregateDeletedEvent)
            aggregate.Delete(e.Data);
    }
    return aggregate;
}
```

**TO BE
CONSIDERED**

Replay is not
about repeating
commands that
generated
events.

**ABOUT
EVENT REPLAY**

Custom Projections of Event Data

Effects may
be redefined
in each
application.

Effects of
occurred
events
applied to
fresh
instances.



Replay just **these** events

Perform different calculations

Apply different forms of business logic

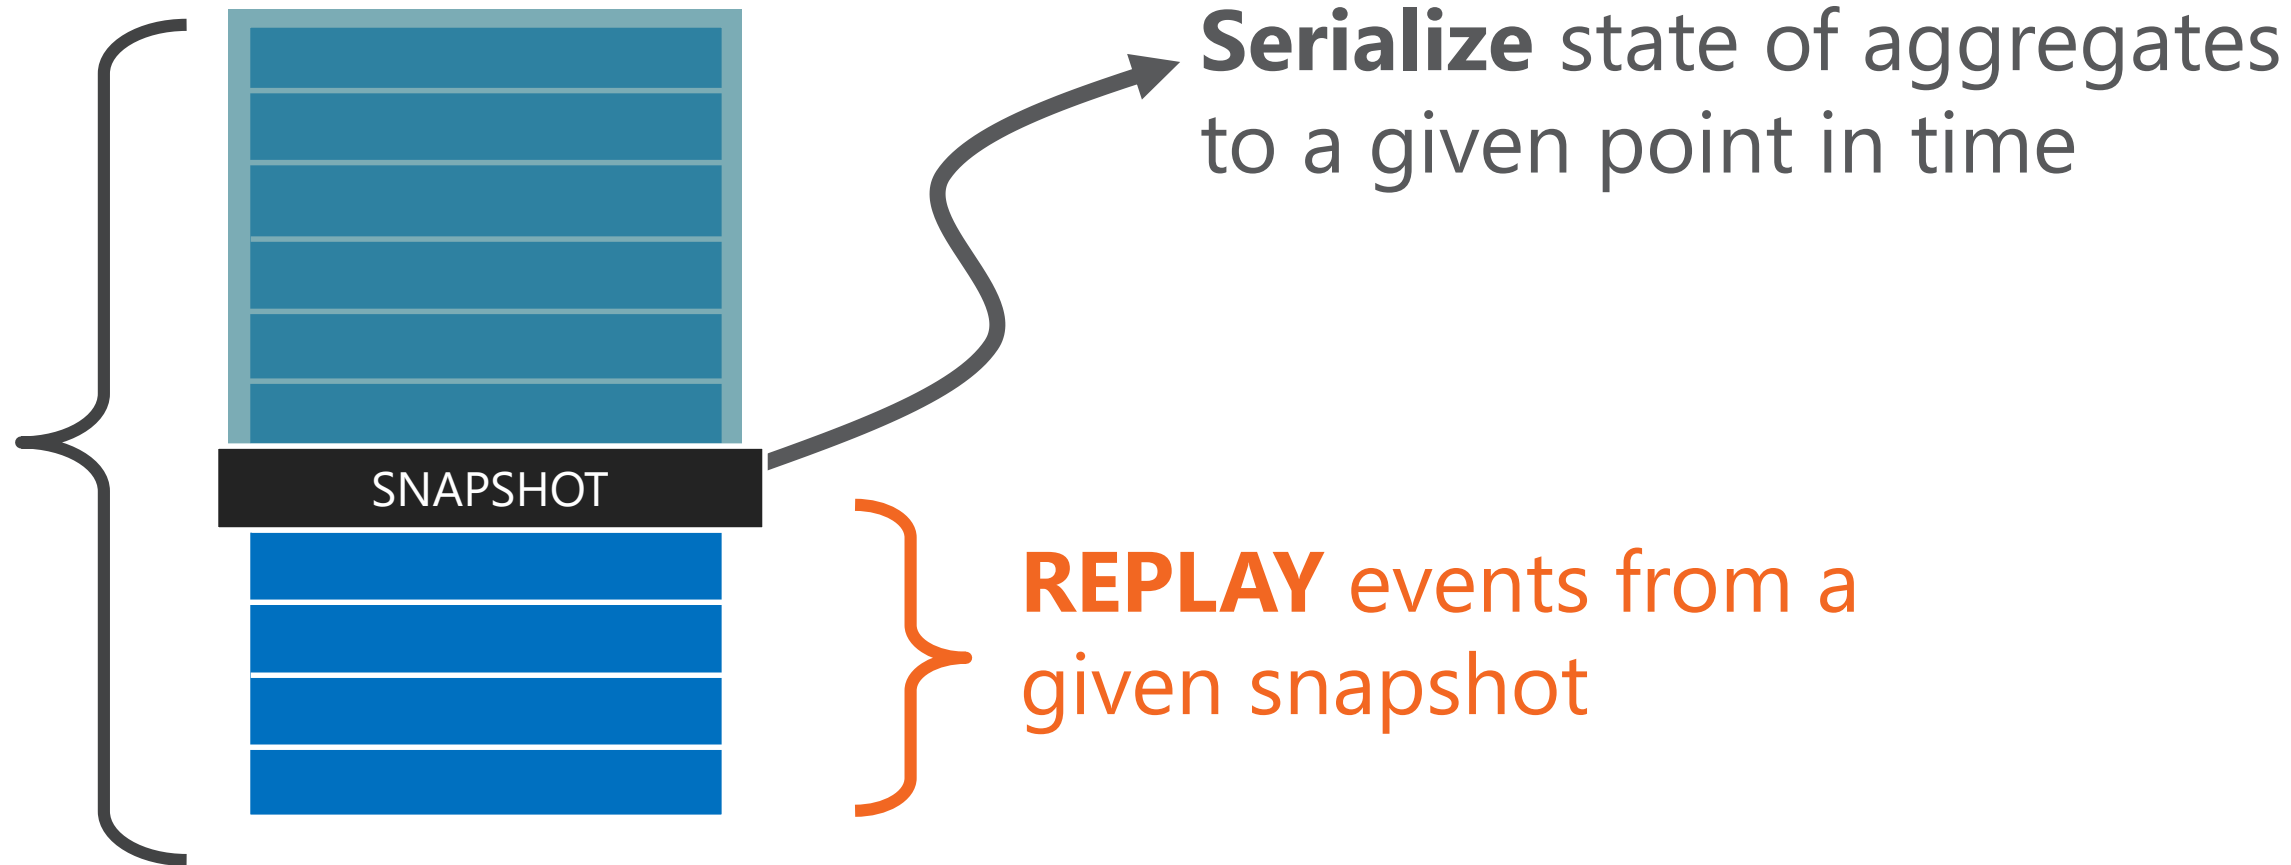
SCALA

Effects of **immutability**

ESSENTIAL



Should you have **performance** concerns?



You can definitely arrange an event sourcing solution all by yourself.

But to store events effectively ad hoc tools may be better.

Event-based Data Stores

Event Store

<http://geteventstore.com>

API

Stream

Stream

Stream

Stream

Stream

Event Store Operations

Writing events to a stream

Reading events from a stream

Subscribing to stream to get updates

```
[
  {
    "eventId": "fbf4a1a1-b4a3-4dfe-a01f-ec52c34e16e4",
    "eventType": "OrderCreatedEvent",
    "data": { "orderId": "1", ... }
  }
]
```

Event Store Types of Subscriptions

Call back a function whenever an event is written to a given stream until the subscription is stopped

Call back a function from a given position up to the end and then turns into volatile.

Multiple consumers are guaranteed to receive at least one notification of events written possibly more.

Volatile

Catch-up

Persistent